

# ECE 153A Lab 1A: Timing Analysis

Luke Lewis | Eduardo Olmos

## Purpose

In this introductory lab for the course, the objective was to familiarize ourselves with the developmental process for the Artix-100T development board. This involved modeling the hardware design of a microblaze processor interconnected with a few peripherals in the Vivado design suite from Xilinx. We then utilized the Xilinx SDK to write C code that runs on our board. Particularly, this lab involved timing various operations in terms of clock cycles and statistically analyze it with consideration that embedded systems are targeted at tightly constrained environments.

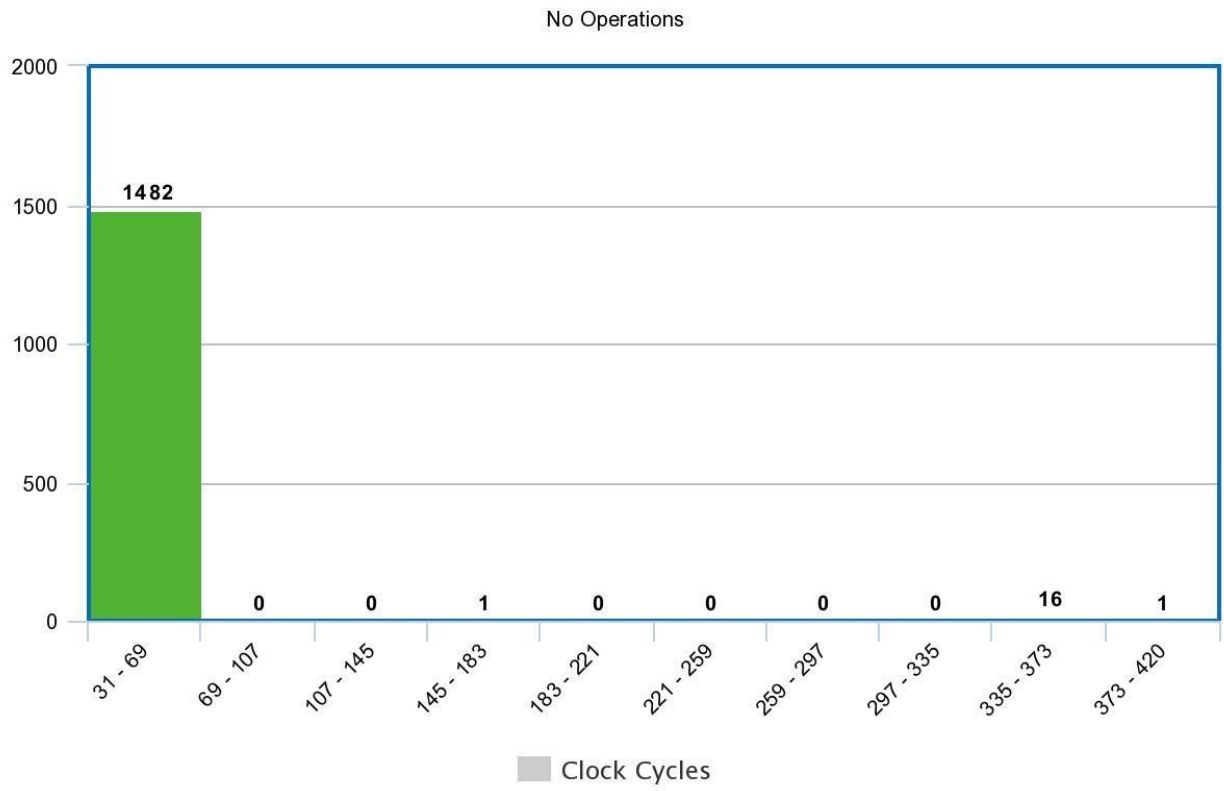
## Operations Measured

- 1) Writing a floating point number using *printf()*
- 2) Writing a string of 10 characters using *xil\_printf()*
- 3) Addition using integers
- 4) Multiplication using integers
- 5) Addition using floating point
- 6) Multiplication using floating point
- 7) Turning on LED
- 8) Turning off LED
- 9) Reading a byte from the DDR2 memory
- 10) Reading 8 bytes from the DDR2 memory


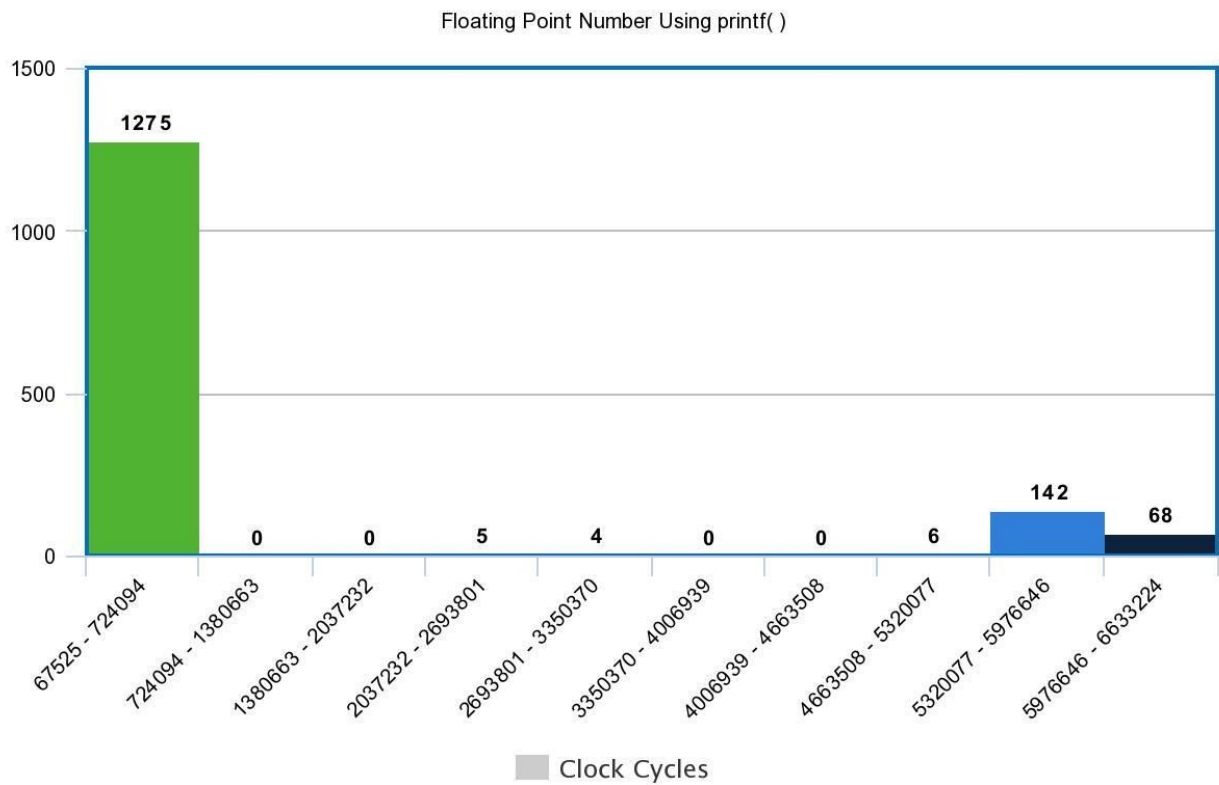
## Assumptions

The way we will be measuring clock cycles will be through the use of a timer. Essentially, for every trial with the operation we want to measure we will have to stop the counter timer (if it is already running), enable the timer, perform the operation, and then record the value. This process loops for the number of trials we want to perform. Before we can start taking measurements, we must account for the activity for when no operations are running. As you can see by observing the histogram below, the processor does not sit idle when given no operation. As detailed in the lab, another method called `extra_method()` is called before we start measurement trials. This method enables an interrupt timer to be handled and is something we need to take into consideration. Measurements in the histogram below show that for 98.8% of the 1500 trials of no operations, clock cycles were in the range of 31-69 and 1.13% of outliers accounting for a greater than tenfold increase in cycles. This means that we also cannot


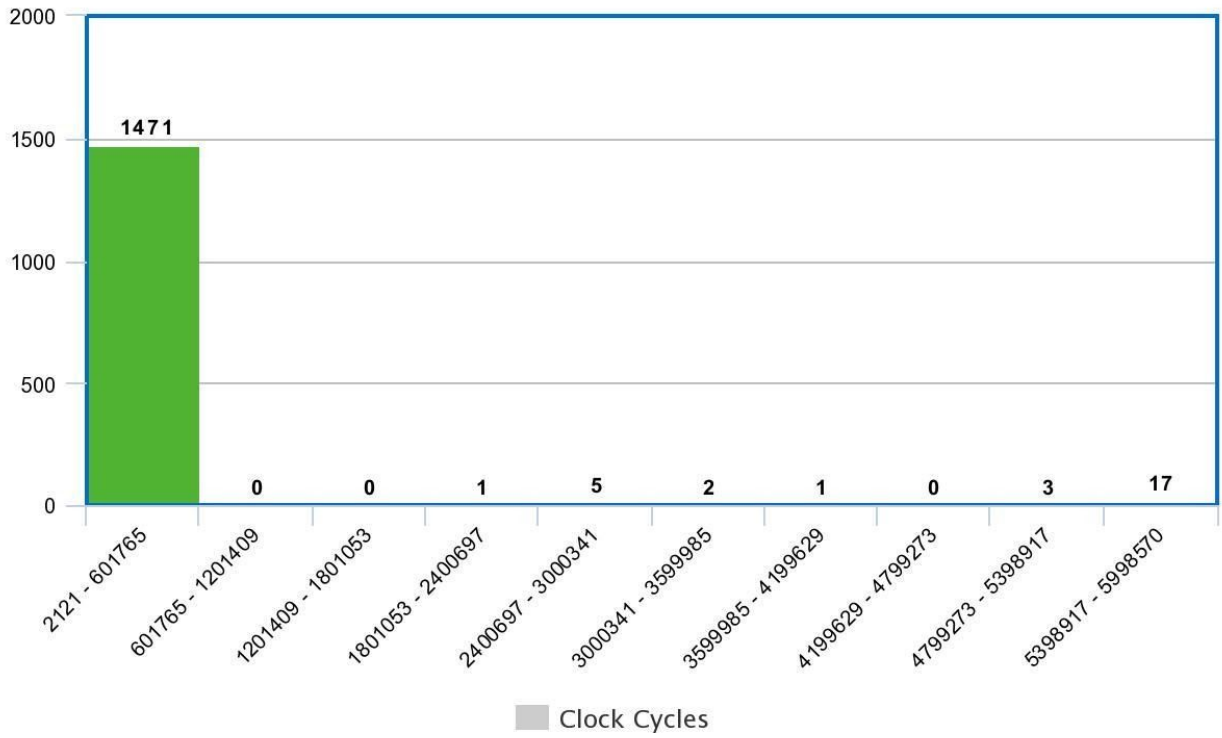
assume that the processor is consistent with each trial of operation or non operation.


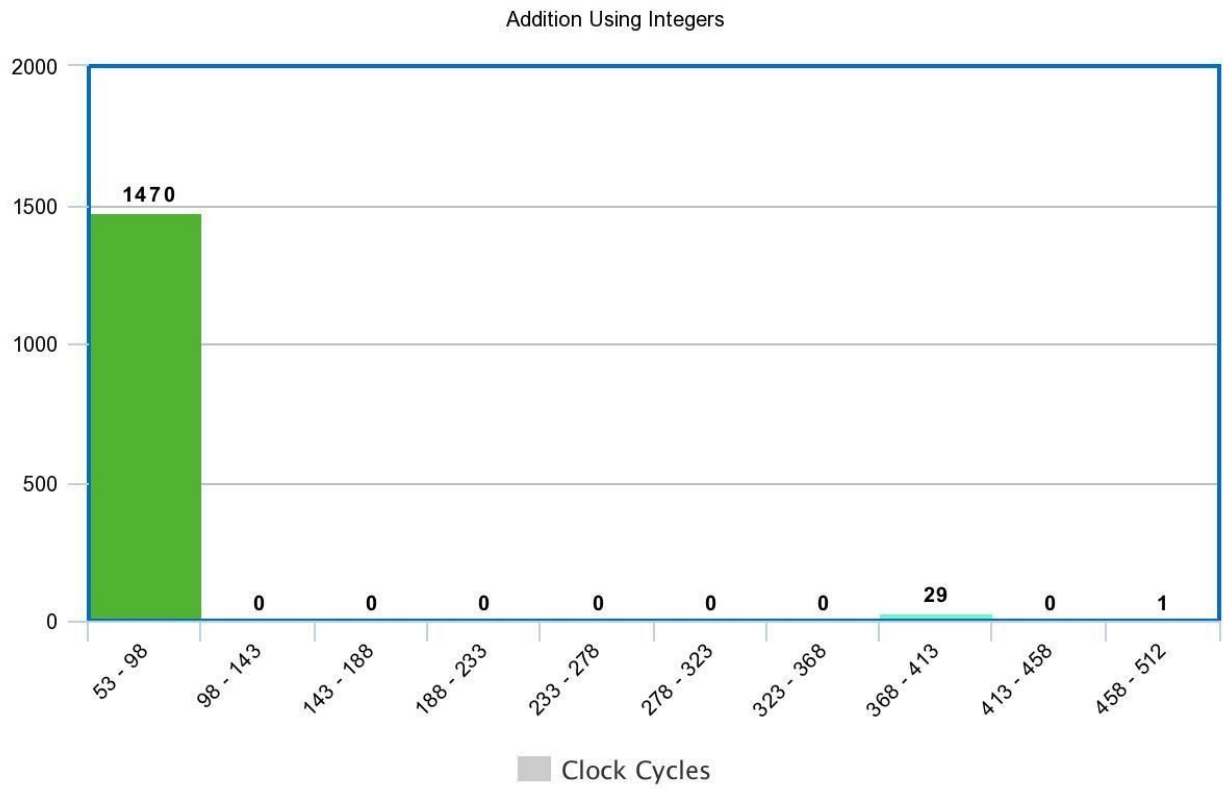



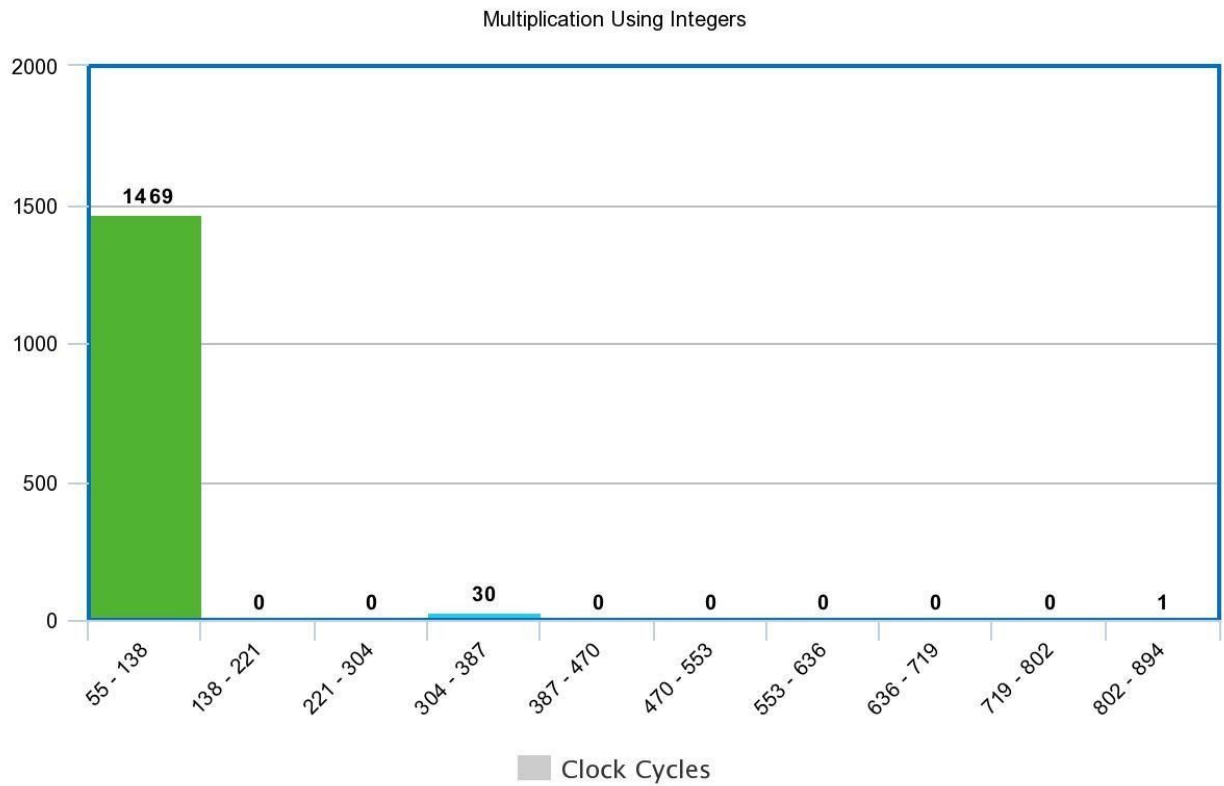
## Results:





String of 10 Characters Using xil\_printf( )


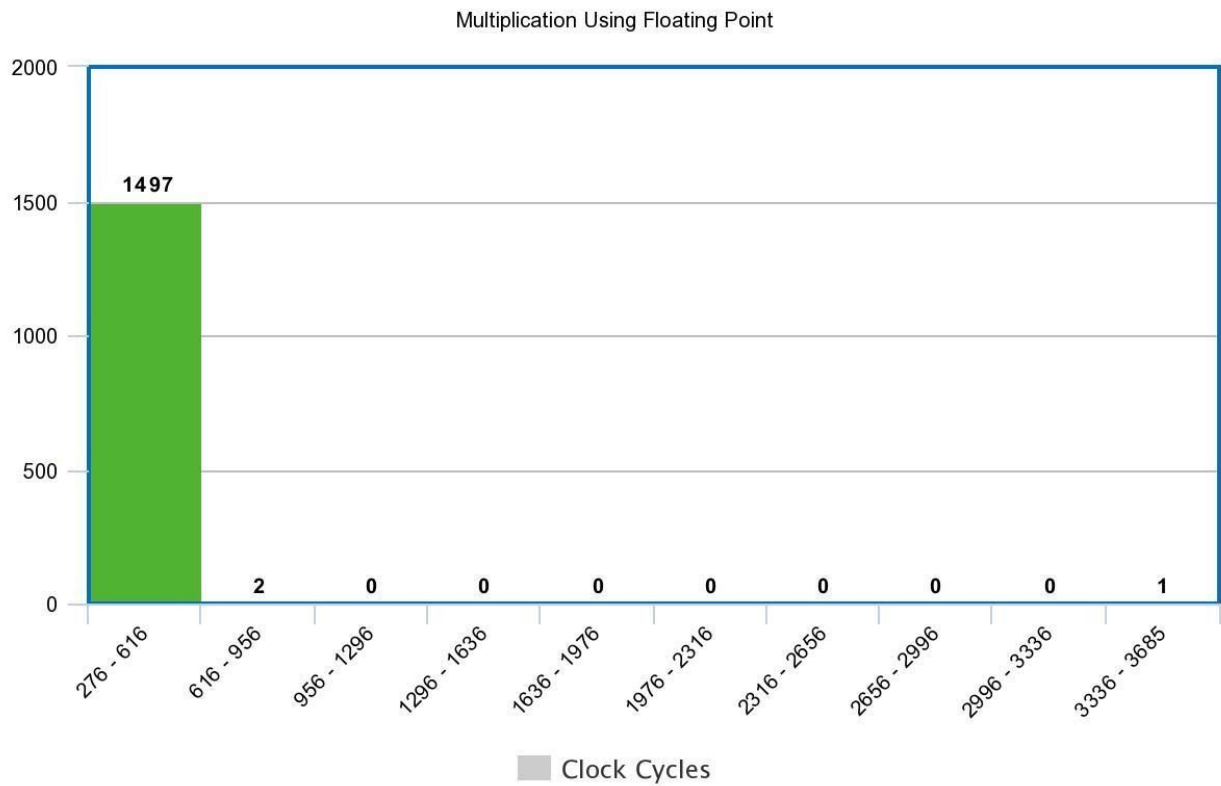










Addition Using Floating Point


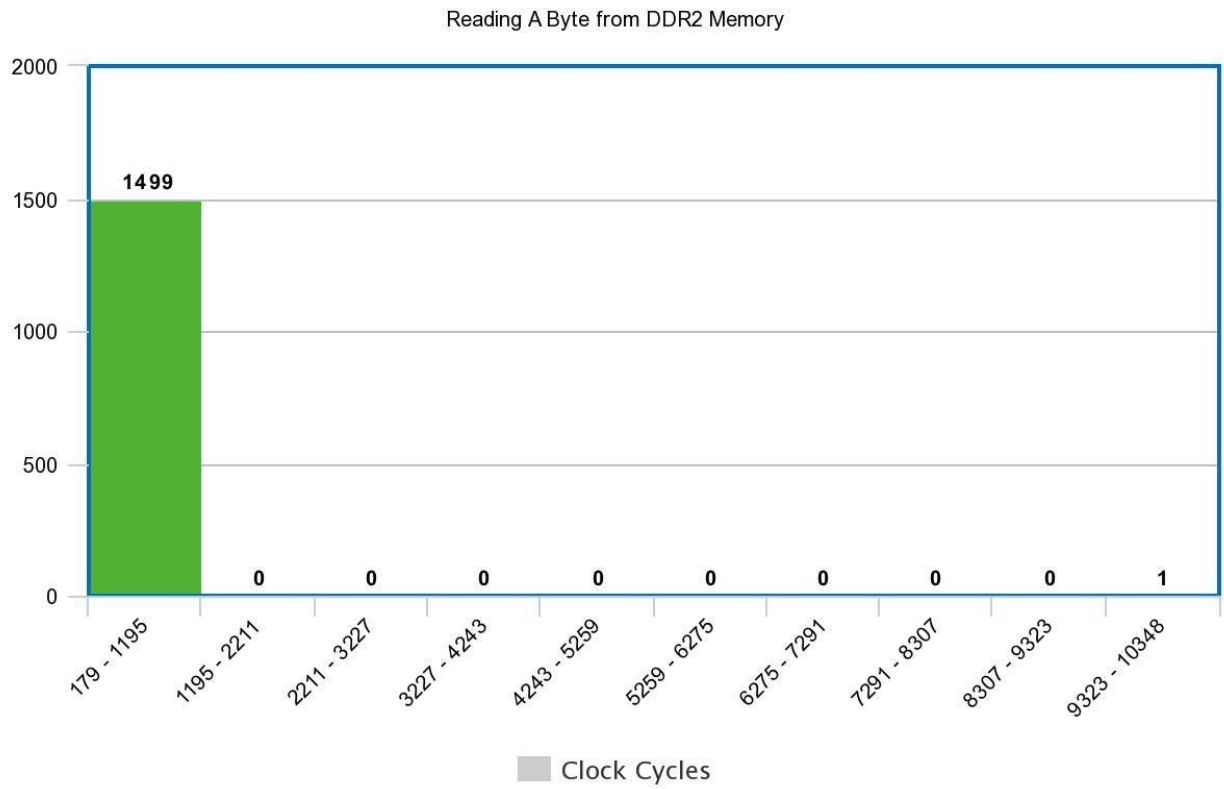



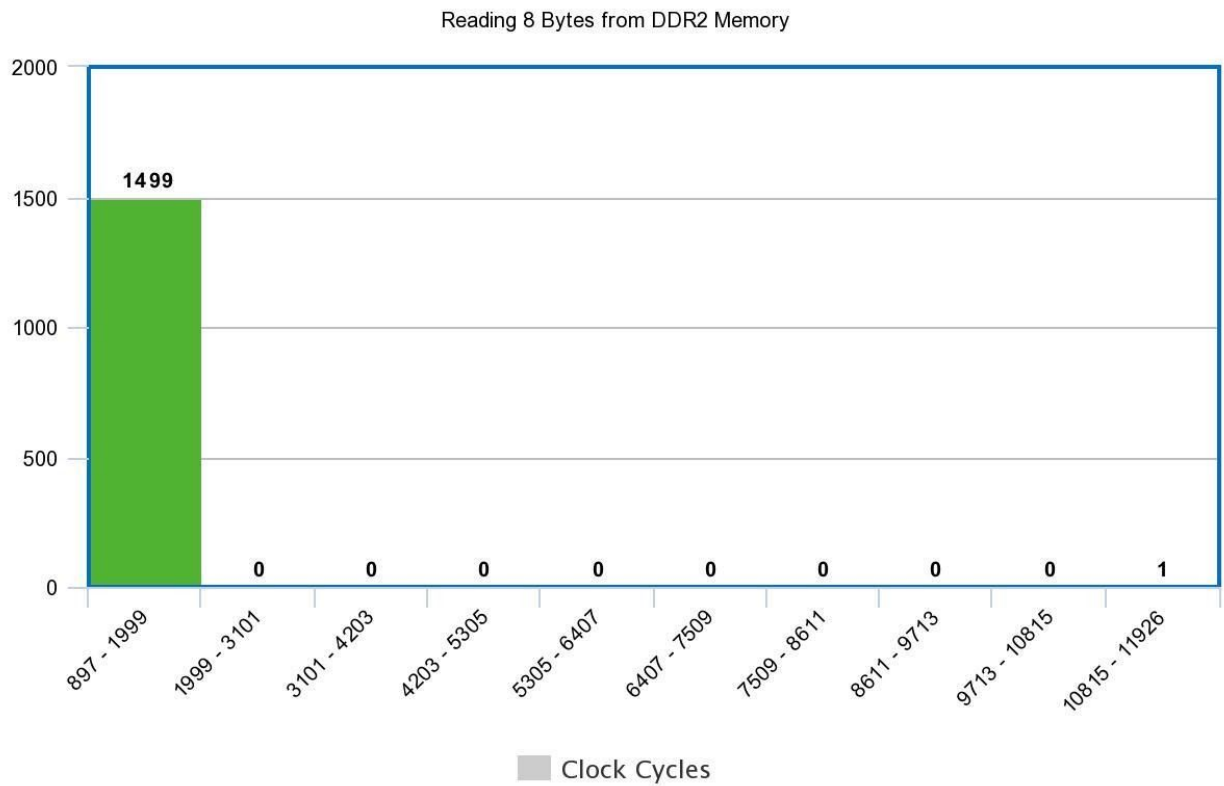












Operation	Min	Max	Expected Value	Inconsistency from Min $\frac{\text{Expected Value} - \text{Min}}{\text{Min}} \times 100$
No Operation	31	411	34.873	12.49%
Printf( ) Float	67525	6633217	918792.438	1260.67%
xil_printf( ) String of 10 Characters	2121	5998566	105817.898	4889.06%
Addition Using Integers	53	508	59.683	12.61%
Multiplication Using Integers	55	889	62.156	13.01%
Addition Using Float	249	3382	281.431	13.02%
Multiplication Using Float	276	3678	311.585	12.89%
Turning LED On	139	1682	156.815	12.82%
Turning LED Off	139	2054	156.815	12.82%
Reading 1 Byte From DDR2 Memory	179	10343	232.784	30.05%
Reading 8 Bytes From DDR2 Memory	897	11923	1168.307	30.25%

## Analysis

In looking at the data we collected, there's a few observations that can be made. Since there is a disparity and spread of values with each operation, we need to take a look at the expected values of these operations as well. The fastest operations, which we will define as the operations with the smallest minimum values of clock cycles, was addition and multiplication using integers. These operations are cost effective relatively compared to the rest of the operations. They are also very similar to each other, with the exception being the 75% greater maximum cycles measurement found with multiplication. These operations are very well implemented and commonly optimized for in hardware so it is no surprise that they are close to the no operation measurements. The expected values found for both operations tell us that they lean close to the minimum values similar to no operation. As a way to measure this distance of the expected value relative to the min, I define the inconsistency value as shown in the chart. We can see that the inconsistency values of addition and multiplication are close to the performance of no operation.

Interestingly, the next quickest group of operations are turning the LED on and off. This involves initializing the GPIO instance and writing on or off at every trial. In looking at the

expected values, they are very close to their minimum value, which are the same for both of them. This is likely due to having a very similar physical procedure in enabling and setting voltages at the GPIO pin. We weren't really sure why the maximum value was higher to turn the led off by 22%. Comparing the inconsistency values, they are actually lower than multiplication using integers. You may say that averaged out these two operations are quick more consistently because their expected values are closer to their minimum.

The next observation we made was that reading a byte from DDR2 memory is faster than addition and multiplication using floating point numbers. While this seems to go against our previously specified point of optimized hardware for addition and multiplication, it seems the more precision that we require due to floating points, an increase in cycle time is to be expected. We do see expected values close to their minimums which means that there must not be much complexity happening.

Reading a byte from DDR2 memory does however have an inconsistency score of 30.05% which means that its expected value is to be found further away than its minimum value. This is reflected in its quite high expected value which is 5778% higher than its minimum value. Looking at the minimum value for an 8 byte read from DDR2 tells us that there is not exactly 8 times the amount of work to be done compared to reading a single byte but it is certainly higher. On the other side of the spectrum, its maximum value is really close to a reading a single byte, so there must be something substantial that affects both operations equally.

Lastly, `printf( )` and `xil_printf( )` were the most expensive operations giving reason to the cautionary tale of bloated printing timing delays in embedded system courses. `xil_printf( )` has the smaller minimum value, it has nearly the same maximum value of `printf( )`. The expected value of `printf( )` is higher than `xil_printf( )` but don't let the inconsistency value fool you. While it is lower than `xil_printf( )`, it is relative to its own minimum value, which is a lot higher than `xil_printf( )`. This is perhaps a failing point of this measurement unit we came up with, but it does allow us to highlight the fact that the expected value of `xil_printf( )` is really far away from its minimum. We also must say that they were both printing different data types and amounts, so that is also a factor.

In conclusion, the data gathered in our trials gave valuable insight on the different amount of cycles that different operations might take, which we can definitely take into account and apply when we create tasks for the processor. The large number of trials, as well as taking measurements of the min, max, expected values, as well as the inconsistency from the min is important in giving a clearer picture of the processor, and a clear picture of expected behaviour through gathered data is absolutely necessary when designing embedded systems, as failure can be very costly and possibly dangerous.